

1a)

Unmodified source code:

- a. Optimization – O3
Output sum: 69999988
Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	20.26	20.13	20.47	19.80	24.19	20.97	19.80
100,000,000	207.12	200.54	199.86	200.88	201.48	201.97	199.86

The shortest run time **for O3 with N = 10,000,000 is 19.80**

The average run time **for O3 with N = 10,000,000 is 20.97**

The shortest run time **for O3 with N = 100,000,000 is 199.86**

The average run time **for O3 with N = 100,000,000 is 201**

- b. Optimization – O2
Output sum: 69999988
Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	19.61	20.45	20.37	23.27	20.20	20.78	19.61
100,000,000	199.00	198.55	198.77	210.63	204.78	202.34	198.55

The shortest run time **for O2 with N = 10,000,000 is 19.61**

The average run time **for O2 with N = 10,000,000 is 20.78**

The shortest run time **for O2 with N = 100,000,000 is 198.55**

The average run time **for O2 with N = 100,000,000 is 202.34**

1b)

Processor Architecture – x86 (64-bits)
CPU Frequency: 1.8 GHz
OS type: macOS Mojave
Standalone system

1c)

Optimization I

Loop Fusion

Resulting Code:

```
Loop Fusion

void do_loops(int *a, int *b, int *c, int N)
{
    int i;
    for (i=N-1; i>=1; i--) {
        a[i] = a[i] + 1;
    }
    for (i=1; i<N; i++) {
        b[i] = a[i+1] + 3;
        c[i] = b[i-1] + 2;
    }
}
```

Optimization – O3

Output sum: 69999988

Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	23.66	24.31	24.18	24.53	26.08	24.55	23.66
100,000,000	586.62	268.78	250.74	253.58	250.05	321.95	250.05

Optimization – O2

Output sum: 69999988

Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	32.75	43.28	33.38	25.91	39.06	34.876	25.91
100,000,000	260.04	255.84	255.38	249.36	253.94	254.91	249.36

The optimization didn't improve the performance of our code. The performance worsened compared to the original unmodified source code. Loop fusion may not improve performance as two separate loops might perform better on an architecture because of data locality.

Optimization II

Loop Peeling

Resulting Code:

```

Loop Peeling

void do_loops(int *a, int *b, int *c, int N)
{
    int i;

    for (i=N-1; i>=1; i-=1) {
        a[i] = a[i] + 1;
    }

    if(N > 0){
        b[1] = a[2] + 3;
    }

    for (i=2; i<N; i++) {
        b[i] = a[i+1] + 3;
    }

    if(N > 0){
        c[1] = b[0] + 3;
    }
    for (i=2; i<N; i++) {
        c[i] = b[i-1] + 2;
    }
}

```

Optimization – O3
 Output sum: 69999988
 Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	28.28	29.20	21.35	21.36	27.41	25.52	21.35
100,000,000	242.08	202.04	201.59	201.11	201.97	209.75	201.11

Optimization – O2
 Output sum: 69999988
 Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	21.10	29.86	21.10	20.90	22.13	23.01	20.90
100,000,000	223.03	201.55	200.97	202.64	200.11	205.66	200.11

The optimization didn't improve the performance of our code. The performance of the modified code is comparable to the unmodified source code. However, the performance of the unmodified source code is better by a few milliseconds.

Optimization III

Loop Reversal

Resulting Code:

```
Loop Reversal

void do_loops(int *a, int *b, int *c, int N)
{
    int i;

    for (i=1; i <= N-1; i+=1) {
        a[i] = a[i] + 1;
    }

    for (i=N-1; i>=1; i--) {
        b[i] = a[i+1] + 3;
    }

    for (i=N-1; i>=1; i--) {
        c[i] = b[i-1] + 2;
    }
}
```

Optimization – O3

Output sum: 69999988

Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	21.04	22.98	27.79	20.97	33.52	25.26	21.04
100,000,000	205.24	195.76	195.95	198.53	199.13	198.92	195.76

Optimization – O2

Output sum: 69999988

Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	21.26	22.22	20.69	23.25	20.94	21.67	20.69
100,000,000	230.76	200.53	200.05	203.38	232.41	213.42	200.05

The performance was improved for the modified code. For optimization level O3 and N =100,000,000, the shortest run time was 195.76 ms which was faster than the shortest run time for the original code which was 199.86 ms. The average run time was also improved for

this case. For optimization level O3 and N = 10,0000, the performance of the modified and the unmodified code was comparable. For optimization level O2, the performance was almost similar to the unmodified code. Loop reversal can be beneficial because some ISAs contain efficient loop count instructions that count in a single direction.

Optimization IV Loop Strip Mining Resulting Code:

```
Loop Strip Mining
void do_loops(int *a, int *b, int *c, int N)
{
    int i;
    int j;

    for (i=N-1; i>=1; i--) {
        a[i] = a[i] + 1;
    }

    for (j=1; j<N; j+=100) {
        for (i=j; i < j+100; i++){
            b[i] = a[i+1] + 3;
        }
    }

    for(j=1; j<N; j+=100){
        for (i=j; i< j+100; i++) {
            c[i] = b[i-1] + 2;
        }
    }
}
```

Optimization – O3
Output sum: 69999988
Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	21.33	26.61	21.17	21.23	21.28	22.32	21.23
100,000,000	619.60	255.99	272.17	245.81	235.14	325.74	235.14

Optimization – O2
Output sum: 69999988
Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	21.49	29.31	21.17	21.47	27.47	24.18	21.17
100,000,000	203.97	210.91	201.46	201.32	201.92	203.91	201.32

The optimization didn't improve the performance of our code. The performance of the modified code was significantly worse for O3 optimization with N = 100,000,000. The performance for O2 optimization was comparable to the unmodified code. Vectorization of the loop didn't yield significant benefit because of processor characteristics.

Optimization V

Loop Unrolling

Resulting Code:

```

Loop Unrolling
void do_loops(int *a, int *b, int *c, int N)
{
    int i;
    for (i=N-1; i>=1; i--) {
        a[i] = a[i] + 1;
    }

    for (i=1; i<N; i+=4) {
        b[i] = a[i+1] + 3;
        b[i+1] = a[i+2] + 3;
        b[i+2] = a[i+3] + 3;
        b[i+3] = a[i+4] + 3;
    }

    for (i=1; i<N; i+=4) {
        c[i] = b[i-1] + 2;
        c[i+1] = b[i] + 2;
        c[i+2] = b[i+1] + 2;
        c[i+3] = b[i+2] + 2;
    }
}

```

Optimization – O3
Output sum: 69999988
Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	30.15	26.45	29.52	32.34	36.93	31.07	26.45
100,000,000	233.25	213.10	207.22	210.03	207.73	214.26	207.22

Optimization – O2
Output sum: 69999988
Runtime is in milliseconds

N	Run 1	Run 2	Run 3	Run 4	Run 5	Average	Shortest
10,000,000	31.32	29.19	28.56	26.45	32.03	29.51	26.45

100,000,000	1148.70	392.75	304.56	215.83	537.39	519.84	215.83
-------------	---------	--------	--------	--------	--------	--------	--------

The optimization didn't improve the performance of our code. The performance was significantly worse for the case of O2 optimization with $N = 100,000,000$. For, all the other cases, the unmodified code performed better than the modified code.

Assembly Code analysis:

a. Unmodified Code:

```

00000000100000c06    subl    $0x1, %ecx
00000000100000c09    movl    %ecx, -0x20(%rbp)
00000000100000c0c    cmpl    $0x1, -0x20(%rbp)
00000000100000c10    jl      0x100000c3d
00000000100000c16    movq    -0x8(%rbp), %rax
00000000100000c1a    movslq   -0x20(%rbp), %rcx
00000000100000c1e    movl    (%rax,%rcx,4), %edx
00000000100000c21    addl    $0x1, %edx
00000000100000c24    movq    -0x8(%rbp), %rax
00000000100000c28    movslq   -0x20(%rbp), %rcx
00000000100000c2c    movl    %edx, (%rax,%rcx,4)
00000000100000c2f    movl    -0x20(%rbp), %eax
00000000100000c32    addl    $-0x1, %eax
00000000100000c35    movl    %eax, -0x20(%rbp)
00000000100000c38    jmp     0x100000c0c
00000000100000c3d    movl    $0x1, -0x20(%rbp)
00000000100000c44    movl    -0x20(%rbp), %eax
00000000100000c47    cmpl    -0x1c(%rbp), %eax
00000000100000c4a    jge     0x100000c7c
00000000100000c50    movq    -0x8(%rbp), %rax
00000000100000c54    movl    -0x20(%rbp), %ecx
00000000100000c57    addl    $0x1, %ecx
00000000100000c5a    movslq   %ecx, %rdx
00000000100000c5d    movl    (%rax,%rdx,4), %ecx
00000000100000c60    addl    $0x3, %ecx
00000000100000c63    movq    -0x10(%rbp), %rax
00000000100000c67    movslq   -0x20(%rbp), %rdx
00000000100000c6b    movl    %ecx, (%rax,%rdx,4)
00000000100000c6e    movl    -0x20(%rbp), %eax
00000000100000c71    addl    $0x1, %eax
00000000100000c74    movl    %eax, -0x20(%rbp)
00000000100000c77    jmp     0x100000c44
00000000100000c7c    movl    $0x1, -0x20(%rbp)
00000000100000c83    movl    -0x20(%rbp), %eax
00000000100000c86    cmpl    -0x1c(%rbp), %eax
00000000100000c89    jge     0x100000cbb
00000000100000c8f    movq    -0x10(%rbp), %rax

```

In the unmodified code, we observe that three loops are being executed with the help of jump and compare statements

b. Loop Fusion:

```
_do_loops:
0000000100000c10    pushq    %rbp
0000000100000c11    movq     %rsp, %rbp
0000000100000c14    movq     %rdi, -0x8(%rbp)
0000000100000c18    movq     %rsi, -0x10(%rbp)
0000000100000c1c    movq     %rdx, -0x18(%rbp)
0000000100000c20    movl     %ecx, -0x1c(%rbp)
0000000100000c23    movl     -0x1c(%rbp), %ecx
0000000100000c26    subl     $0x1, %ecx
0000000100000c29    movl     %ecx, -0x20(%rbp)
0000000100000c2c    cmpl     $0x1, -0x20(%rbp)
0000000100000c30    jnl      0x100000c5d
0000000100000c36    movq     -0x8(%rbp), %rax
0000000100000c3a    movslq   -0x20(%rbp), %rcx
0000000100000c3e    movl     (%rax,%rcx,4), %edx
0000000100000c41    addl     $0x1, %edx
0000000100000c44    movq     -0x8(%rbp), %rax
0000000100000c48    movslq   -0x20(%rbp), %rcx
0000000100000c4c    movl     %edx, (%rax,%rcx,4)
0000000100000c4f    movl     -0x20(%rbp), %eax
0000000100000c52    addl     $-0x1, %eax
0000000100000c55    movl     %eax, -0x20(%rbp)
0000000100000c58    jmp      0x100000c2c
0000000100000c5d    movl     $0x1, -0x20(%rbp)
0000000100000c64    movl     -0x20(%rbp), %eax
0000000100000c67    cmpl     -0x1c(%rbp), %eax
0000000100000c6a    jge      0x100000cba
0000000100000c70    movq     -0x8(%rbp), %rax
0000000100000c74    movl     -0x20(%rbp), %ecx
0000000100000c77    addl     $0x1, %ecx
0000000100000c7a    movslq   %ecx, %rdx
0000000100000c7d    movl     (%rax,%rdx,4), %ecx
0000000100000c80    addl     $0x3, %ecx
0000000100000c83    movq     -0x10(%rbp), %rax
0000000100000c87    movslq   -0x20(%rbp), %rdx
0000000100000c8b    movl     %ecx, (%rax,%rdx,4)
0000000100000c8e    movq     -0x10(%rbp), %rax
```

In the loop-fusion modification, we observe that two loops are being executed with the help of jump and compare statements.

c. Loop Peeling:

```

0000000100000c07    movq    -0x8(%rbp), %rax
0000000100000c0b    movl    0x8(%rax), %ecx
0000000100000c0e    addl    $0x3, %ecx
0000000100000c11    movq    -0x10(%rbp), %rax
0000000100000c15    movl    %ecx, 0x4(%rax)
0000000100000c18    movl    $0x2, -0x20(%rbp)
0000000100000c1f    movl    -0x20(%rbp), %eax
0000000100000c22    cmpl    -0x1c(%rbp), %eax
0000000100000c25    jge     0x100000c57
0000000100000c2b    movq    -0x8(%rbp), %rax
0000000100000c2f    movl    -0x20(%rbp), %ecx
0000000100000c32    addl    $0x1, %ecx
0000000100000c35    movslq  %ecx, %rdx
0000000100000c38    movl    (%rax,%rdx,4), %ecx
0000000100000c3b    addl    $0x3, %ecx
0000000100000c3e    movq    -0x10(%rbp), %rax
0000000100000c42    movslq  -0x20(%rbp), %rdx
0000000100000c46    movl    %ecx, (%rax,%rdx,4)

```

For the loop peeling example, we observe the addition of 3 to the first element of the arrays, before the loop initiates.

d. Loop Reversal:

```

0000000100000bfa    movl    -0x20(%rbp), %eax
0000000100000bfd    movl    -0x1c(%rbp), %ecx
0000000100000c00    subl    $0x1, %ecx
0000000100000c03    cmpl    %ecx, %eax
0000000100000c05    jg      0x100000c32
0000000100000c0b    movq    -0x8(%rbp), %rax
0000000100000c0f    movslq  -0x20(%rbp), %rcx
0000000100000c13    movl    (%rax,%rcx,4), %edx
0000000100000c16    addl    $0x1, %edx
0000000100000c19    movq    -0x8(%rbp), %rax
0000000100000c1d    movslq  -0x20(%rbp), %rcx
0000000100000c21    movl    %edx, (%rax,%rcx,4)
0000000100000c24    movl    -0x20(%rbp), %eax
0000000100000c27    addl    $0x1, %eax
0000000100000c2a    movl    %eax, -0x20(%rbp)
0000000100000c2d    jmp     0x100000bfa

```

For loop reversal, we observe that the code is similar to the unmodified code, however, the order of loop has been reversed and the loop counter is decremented with the help of the subl instruction.

e.

Loop Strip-Mining:

```
00000000100000bca    movsq    -0x20(%rbp), %ecx
00000000100000bcc    movl     %edx, (%rax,%rcx,4)
00000000100000bcf    movl     -0x20(%rbp), %eax
00000000100000bd2    addl     $-0x1, %eax
00000000100000bd5    movl     %eax, -0x20(%rbp)
00000000100000bd8    jmp      0x100000bac
00000000100000bdd    movl     $0x1, -0x24(%rbp)
00000000100000be4    movl     -0x24(%rbp), %eax
00000000100000be7    cmpl     -0x1c(%rbp), %eax
00000000100000bea    jge      0x100000c46
00000000100000bf0    movl     -0x24(%rbp), %eax
00000000100000bf3    movl     %eax, -0x20(%rbp)
00000000100000bf6    movl     -0x20(%rbp), %eax
00000000100000bf9    movl     -0x24(%rbp), %ecx
00000000100000bfc    addl     $0x64, %ecx
00000000100000bff    cmpl     %ecx, %eax
00000000100000c01    jge      0x100000c33
00000000100000c07    movq     -0x8(%rbp), %rax
00000000100000c0b    movl     -0x20(%rbp), %ecx
00000000100000c0e    addl     $0x1, %ecx
00000000100000c11    movslq   %ecx, %rdx
00000000100000c14    movl     (%rax,%rdx,4), %ecx
00000000100000c17    addl     $0x3, %ecx
00000000100000c1a    movq     -0x10(%rbp), %rax
00000000100000c1e    movslq   -0x20(%rbp), %rdx
00000000100000c22    movl     %ecx, (%rax,%rdx,4)
00000000100000c25    movl     -0x20(%rbp), %eax
00000000100000c28    addl     $0x1, %ecx
```

For loop strip-mining, we observe a double nested loop structure. The number of instructions also increased significantly.

f. Partial Loop Unrolling:

```

0000000100000b90      addl    $0x3, %ecx
0000000100000b93      movq    -0x10(%rbp), %rax
0000000100000b97      movslq  -0x20(%rbp), %rdx
0000000100000b9b      movl    %ecx, (%rax,%rdx,4)
0000000100000b9e      movq    -0x8(%rbp), %rax
0000000100000ba2      movl    -0x20(%rbp), %ecx
0000000100000ba5      addl    $0x2, %ecx
0000000100000ba8      movslq  %ecx, %rdx
0000000100000bab      movl    (%rax,%rdx,4), %ecx
0000000100000bae      addl    $0x3, %ecx
0000000100000bb1      movq    -0x10(%rbp), %rax
0000000100000bb5      movl    -0x20(%rbp), %esi
0000000100000bb8      addl    $0x1, %esi
0000000100000bbb      movslq  %esi, %rdx
0000000100000bbe      movl    %ecx, (%rax,%rdx,4)
0000000100000bc1      movq    -0x8(%rbp), %rax
0000000100000bc5      movl    -0x20(%rbp), %ecx
0000000100000bc8      addl    $0x3, %ecx
0000000100000bcb      movslq  %ecx, %rdx
0000000100000bce      movl    (%rax,%rdx,4), %ecx
0000000100000bd1      addl    $0x3, %ecx
0000000100000bd4      movq    -0x10(%rbp), %rax
0000000100000bd8      movl    -0x20(%rbp), %esi
0000000100000bdb      addl    $0x2, %esi
0000000100000bde      movslq  %esi, %rdx
0000000100000be1      movl    %ecx, (%rax,%rdx,4)
0000000100000be4      movq    -0x8(%rbp), %rax
0000000100000be8      movl    -0x20(%rbp), %ecx
0000000100000beb      addl    $0x4, %ecx
0000000100000bec      movslq  %ecx, %rdx

```

For the partial loop unrolling, we observe more statements being executed within the loop and the number of instructions increase significantly.

1d)

I could beat the performance of the compiler for only case when I optimized the loop with the help of loop reversal. On an average, the performance of the compiler was significantly better for other scenarios. Compilers are great at knowing the better way to perform an operation or sequence of operations in the context of the target and compilation objectives. Hence, the overall optimization done by the compiler produced better results than my standalone modifications.

5. Loop Transformations

(a) Loop Fusion

Original Dependencies:

S1 => T S2 (loop independent)

S1 => O S3 (loop carried)

S2 => A S3 (loop carried)

After Loop Fusion, the dependencies become as follow:

S1 => T S2(loop independent)

$S1 \Rightarrow O S3(\text{loop carried})$
 $S2 \Rightarrow T S3(\text{loop carried})$

Loop fusion is safe if and only if no data dependence between the nests becomes loop-carried data dependence of a different type. We observe that the condition is violated for the above loop fusion.

(b) Loop Interchange

Original Dependencies:

$S1 \Rightarrow A S1$ (loop dependent)

We observe that there is a data dependence carried by the outer loop executed for i and j to another statement instance i' and j' where $i < i'$ and $j > j'$.

Example: $(1,3) \Rightarrow (2,2)$

Loop Interchange is safe if outermost loop does not carry any data dependence from one statement instance executed for i and j to another statement instance executed for i' and j' where $(i < i' \text{ and } j > j') \text{ OR } (i > i' \text{ and } j < j')$. Since this condition is violated, loop interchange won't be safe

(c) Loop Fission

Original Dependencies:

$S1 \Rightarrow T S2$ (loop dependent)

After loop fission:

$S1 \Rightarrow T S2$ (loop dependent)

Loop Fission is safe if and only if statements involved in a cycle of loop-carried data dependences remain in the same loop and if there exists a data dependence between two statements placed in different loops, the dependence type must not change. There is no violation of condition for this case. Hence, it is safe to use loop fission.

6. Loop Transformation

Initial Code:

```
int a[N][4];
int rand_number = rand();

for (i=0; i<4; i++) {
    threshold = 2.0 * rand_number;

    for (j=0; j<N; j++) {
        if (threshold < 4) {
            sum = sum + a[j][i];
        }
        else {
            sum = sum + a[j][i] + 1;
        }
    }
}
```

A. Loop Invariant

Pulling non-loop-dependent calculations out of the loop

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;

for (i=0; i<4; i++) {
    for (j=0; j<N; j++) {
        if (threshold < 4) {
            sum = sum + a[j][i];
        }
        else {
            sum = sum + a[j][i] + 1;
        }
    }
}
```

B. Loop un-switching

Move a conditional expression outside of a loop, and replicate loop body inside of each conditional block

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;
for (i=0; i<4; i++) {
    if (threshold < 4) {
        for (j=0; j<N; j++){
            sum = sum + a[j][i];
        }
    }
    else {
        for (j=0; j<N; j++){
            sum = sum + a[j][i] + 1;
        }
    }
}
```

C. Loop Interchange

Switch the positions of one loop that is tightly nested within another loop

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;
for (j=0; j<N; j++) {
    if (threshold < 4) {
        for (i=0; i<4; i++){
            sum = sum + a[j][i];
        }
    }
    else {
        for (i=0; i<4; i++){
            sum = sum + a[j][i] + 1;
        }
    }
}
```

D. Loop Peeling

Remove first and/or last iterations of a loop body to separate code outside the loop

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;
for (j=0; j<N; j++) {
    if (threshold < 4) {
        sum = sum + a[j][0];
        for (i=1; i<4; i++){
            sum = sum + a[j][i];
        }
    }
    else {
        sum = sum + a[j][0] + 1;
        for (i=1; i<4; i++){
            sum = sum + a[j][i] + 1;
        }
    }
}
```

E. Loop Unrolling

Combine multiple instances of the loop body and make corresponding reduction to the loop iteration count

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;
for (j=0; j<N; j++) {
    if (threshold < 4) {
        sum = sum + a[j][0];
        sum = sum + a[j][1];
        sum = sum + a[j][2];
        sum = sum + a[j][3];
    }
    else {
        sum = sum + a[j][0] + 1;
        sum = sum + a[j][1] + 1;
        sum = sum + a[j][2] + 1;
        sum = sum + a[j][3] + 1;
    }
}
```

F. Loop Reversal

Reverse the order of the loop iteration

```
int a[N][4];
int rand_number = rand();
threshold = 2.0 * rand_number;
for (j=N-1; j>=0; j++) {
    if (threshold < 4) {
        sum = sum + a[j][0];
        sum = sum + a[j][1];
        sum = sum + a[j][2];
        sum = sum + a[j][3];
    }
    else {
        sum = sum + a[j][0] + 1;
        sum = sum + a[j][1] + 1;
        sum = sum + a[j][2] + 1;
        sum = sum + a[j][3] + 1;
    }
}
```

3. Function in-lining and performance

(a)

In-lining the add function

```
checksum=-2120047872
Time=101.488 milliseconds
checksum=-2120047872
Time=101.525 milliseconds
checksum=-2120047872
Time=121.764 milliseconds
checksum=-2120047872
Time=104.033 milliseconds
checksum=-2120047872
Time=102.093 milliseconds
```

Average Time: 106.1806 ms

Shortest Run-time: 101.488 ms

Never In-lining the add function

```
checksum=-2120047872
Time=228.806 milliseconds
checksum=-2120047872
Time=211.334 milliseconds
checksum=-2120047872
Time=202.758 milliseconds
checksum=-2120047872
Time=198.533 milliseconds
checksum=-2120047872
Time=202.362 milliseconds
```

Average Time: 208.7586

Shortest Run-time: 202.362 ms

(b) Effects of in-lining:

Inline function is a function that is expanded in line when it is called. When the inline function is called whole code of the inline function gets inserted or substituted at the point of inline function call. This substitution is performed by the C++ compiler at compile time. Inline function may increase efficiency if it is small.

Code-snippets for in-lining:

```
00000001000014fd    callq  0x100001cc0 ## symbol stub for: _gettimeofday
0000000100001502    movq   -0x38(%rbp), %rsi
0000000100001506    movl   %r12d, -0x50(%rbp)
000000010000150a    cmpl   $0x7, %r12d
000000010000150e    jbe    0x100001600
```

```
0000000100001600    movl   (%r15,%rbx,4), %eax
0000000100001604    addl   (%r14,%rbx,4), %eax
0000000100001608    movl   %eax, (%r13,%rbx,4)
000000010000160d    incq   %rbx
0000000100001610    cmpq   %rbx, %rsi
0000000100001613    jne    0x100001600
0000000100001615    xorl   %r12d, %r12d
0000000100001618    leaq   -0x70(%rbp), %rdi
000000010000161c    xorl   %esi, %esi
000000010000161e    callq  0x100001cc0 ## symbol stub for: _gettimeofday
```

The loop being carried out

```
00000000100001290      pushq    %rbp
00000000100001291      movq     %rsp, %rbp
00000000100001294      leal     (%rdi,%rsi), %eax
00000000100001297      popq     %rbp
00000000100001298      retq
00000000100001299      nopl     (%rax)
```

We don't observe any explicit call to the addition function (`_Z3addii`) as it gets expanded within another function

Code snippet-for no in-lining:

```
00000000100001609      callq    0x100001cc0 ## symbol stub for: _gettimeofday
0000000010000160e      nop
00000000100001610      movl     (%r12,%rbx,4), %edi
00000000100001614      movl     (%r15,%rbx,4), %esi
00000000100001618      callq    _Z3addii ## add(int, int)
0000000010000161d      movl     %eax, (%r13,%rbx,4)
00000000100001622      incq     %rbx
00000000100001625      cmpq     %rbx, %r14
00000000100001628      jne      0x100001610
0000000010000162a      xorl     %r12d, %r12d
0000000010000162d      leaq     -0x78(%rbp), %rdi
00000000100001631      xorl     %esi, %esi
00000000100001633      callq    0x100001cc0 ## symbol stub for: _gettimeofday
```

We observe an explicit call to the addition function

- (c) My measured performance results match my expectations. The average runtime of the program with always in-lining attribute was almost approximately half of the average runtime of the program with no in-lining. In-lining reduces the number of executed instructions by avoiding function calls and return instructions whereas context-switch can prove to be costly for small functions. For our program we are calling the addition function several times and since the size of the function is relatively small, we get superior performance with in-lining.

(d) Performance of original code:

```
checksum=-2120047872
Time=101.374 milliseconds
checksum=-2120047872
Time=101.923 milliseconds
checksum=-2120047872
Time=106.374 milliseconds
checksum=-2120047872
Time=108.155 milliseconds
checksum=-2120047872
Time=100.814 milliseconds
```

Average Run-time: 103.728 ms


Shortest Run-time: 101.374 ms

The performance of the original code is almost the same as the performance of the in-lined code. Based on this, we can make the claim that the compiler is in-lining the add function.

2.

(a)

Iteration-space Traversal Graph



Iteration: $(1,1) \Rightarrow (2,1) \Rightarrow (2,2) \Rightarrow (3,1) \Rightarrow (3,2)$

(b)

Dependence:

- | | |
|---|------------------|
| 1. $S1 \Rightarrow T S3$ | Loop Independent |
| 2. $S1 \Rightarrow A S2$ | Loop Independent |
| 3. $S3[i,j] \Rightarrow T S1[i+1,j]$ | Loop Carried |
| 4. $S4[i,j] \Rightarrow T S4[i+1, j-1]$ | Loop Carried |
| 5. $S1[i,j] \Rightarrow T S2[i+1,j+1]$ | Loop Carried |
| 6. $S1[i,j] \Rightarrow A S1[i+1,j-1]$ | Loop Carried |
| 7. $S3[i,j] \Rightarrow T S2[i+2,j]$ | Loop Carried |

(c)

Loop Dependence Traversal Graph

